

HBASE-19876

chia7712

2018-01-28

Contents

Chapter 1

Root issue HBASE-19876

1.1 Summary

The exception happening in converting pb mutation to hbase.mutation messes up the CellScanner

1.2 Description

```
1 2018-01-27 22:51:43,794 INFO [hconnection-0x3291b443-shared-pool11-t6] client.  
   AsyncRequestFutureImpl(778): id=5, table=testQuotaStatusFromMaster3, attempt=6/16  
   failed=20ops, last exception=org.apache.hadoop.hbase.client.WrongRowIOException: org  
   .apache.hadoop.hbase.client.WrongRowIOException: The row in xxx doesn't match the  
   original one aaa  
2   at org.apache.hadoop.hbase.client.Mutation.add(Mutation.java:776)  
3   at org.apache.hadoop.hbase.client.Put.add(Put.java:282)  
4   at org.apache.hadoop.hbase.shaded.protobuf.ProtobufUtil.toPut(ProtobufUtil.java:642)  
5   at org.apache.hadoop.hbase.regionserver.RSRpcServices.doBatchOp(RSRpcServices.java  
   :952)  
6   at org.apache.hadoop.hbase.regionserver.RSRpcServices.doNonAtomicRegionMutation(  
   RSRpcServices.java:896)  
7   at org.apache.hadoop.hbase.regionserver.RSRpcServices.multi(RSRpcServices.java:2591)  
8   at org.apache.hadoop.hbase.shaded.protobuf.generated.ClientProtos$ClientService$.  
   callBlockingMethod(ClientProtos.java:41560)  
9   at org.apache.hadoop.hbase.ipc.RpcServer.call(RpcServer.java:404)  
10  at org.apache.hadoop.hbase.ipc.CallRunner.run(CallRunner.java:130)  
11  at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:324)  
12  at org.apache.hadoop.hbase.ipc.RpcExecutor$Handler.run(RpcExecutor.java:304)
```

I noticed this bug when testing the table space quota.

When rs are converting pb mutation to hbase.mutation, the quota exception or cell exception may be thrown.

```
1 Unable to find source-code formatter for language: rsrpcservices#dobatchop.java.  
   Available languages are: actionscript, ada, applescript, bash, c, c#, c++, cpp, css,  
   erlang, go, groovy, haskell, html, java, javascript, js, json, lua, none, nyan,  
   objc, perl, php, python, r, rainbow, ruby, scala, sh, sql, swift, visualbasic, xml,  
   yaml  
   for (ClientProtos.Action action: mutations) {  
2     MutationProto m = action.getMutation();  
3     Mutation mutation;  
4     if (m.getMutateType() == MutationType.PUT) {  
5       mutation = ProtobufUtil.toPut(m, cells);  
6       batchContainsPuts = true;  
7     } else {  
8       mutation = ProtobufUtil.toDelete(m, cells);
```

```

9         batchContainsDelete = true;
10    }
11    mutationActionMap.put(mutation, action);
12    mArray[i++] = mutation;
13    checkCellSizeLimit(region, mutation);
14    // Check if a space quota disallows this mutation
15    spaceQuotaEnforcement.getPolicyEnforcement(region).check(mutation);
16    quota.addMutation(mutation);
17 }

```

rs has caught the exception but it doesn't have the cellscanner skip the failed cells.

```

1    } catch (IOException ie) {
2        if (atomic) {
3            throw ie;
4        }
5        for (Action mutation : mutations) {
6            builder.setResultOrException(getResultOrException(ie, mutation.getIndex()));
7        }
8    }

```

The bug results in the WrongRowIOException to remaining mutations since they refer to invalid cells.

1.3 Attachments

1. [HBASE-19876.branch-1.2.v0.patch](#)
2. [HBASE-19876.master.001.patch](#)
3. [HBASE-19876.v0.patch](#)
4. [HBASE-19876.v1.patch](#)
5. [HBASE-19876.v2.patch](#)
6. [HBASE-19876.v3.patch](#)
7. [HBASE-19876.v3.patch](#)
8. [HBASE-19876.v3.patch](#)
9. [HBASE-19876.v3.patch](#)
10. [HBASE-19876.v4.patch](#)
11. [HBASE-19876.v5.patch](#)
12. [HBASE-19876.v6.patch](#)

1.4 Comments

1. **anoop.hbase:** A related issue was fixed some time back.. Skipping the cur cell which is problematic. So that might be missing for some places exception handling !!!

2. **chia7712:** {quote}A related issue was fixed some time back.

{quote}

HBASE-15214

{quote}Skipping the cur cell which is problematic. So that might be missing for some places exception handling !!!

{quote}

yep, all append/increment/doBatchOphave similar issue...

3. **chia7712:** The append/increment wont do the batch so only batch of put/Delete has this issue. That makes the fix easy. Will come back soon.
4. **anoop.hbase:** Append and increment Mutation may have N Cells in that (for same row). Even if one cell is malformed, we have to fail the op. Ya that make sense. Need to verify once though. Thanks Chia.
5. **yuzhihong@gmail.com:** Looks good overall

```
1
2 1008          // if the mArray[i] != null, it means the cell scanner has be read.
```

Typo: 'be'

Please remove commented out code in the test.

6. **chia7712:** v1 address [~yuzhihong@gmail.com]'s comment
7. **chia7712:** The failed tests are caused by new timeout rule.HBASE-19911 is dealing with them.

v2 fix the checkstyle warnings.

8. **chia7712:** v3 fix the failed test...
9. **stack:** .001 rebase of [~chia7712] patch.
10. **chia7712:** Thanks [~stack]. I forget to say this issue is blocked byHBASE-19900. IfHBASE-19900 is not resolved, the tests in the patch are unstable because of the region-level exception.
11. **chia7712:** Trying to write the test for branch-1.2 and branch-1.3...
12. **chia7712:** The failed tests pass locally...Try the patch for master. Ping for reviews also.

13. **stack:** Took a look. This is tricky stuff. Thanks for diving in.

This comment is a little confusing:

```
890 // TODO doBatchOp has handled the IOE for all non-atomic operations
```

```
891 // Catching IOE here may confuse readers in the future
```

It is in a catch for an IOE so ... should we remove the catch block this comment is in?

I was trying to see if you could get rid of the cycle looking for Gets in checkAndRow-Mutate now it has been added to doBatchOp but couldn't figure it; an optimization for another time (smile).

On below, I agree.

```
1011 // The atomic ops use the global exception although I feel it is ok to add the ex-  
exception
```

```
1012 // to each action.
```

We need an issue for the below?

```
184 * The invalid cells is in rm. The rm should fail but the subsequent mutations should  
succeed.
```

```
185 * Currently, we have no client api to submit the request consisting of condition-rm  
and mutation.
```

```
186 * Hence, this test build the request manually.
```

We need to make it so it is easy to stand up a RegionServer w/o a Master so you don't have to do the below. This should be easier now... since HBASE-18846. Let me work on it....

```
218 ClientProtos.MultiResponse response = rs.getRSRpcServices().multi(controller, request);
```

Great tests [~chia7712]

This patch is super important stuff.

+1 to commit.

14. **chia7712:** {quote}It is in a catch for an IOE so ... should we remove the catch block this comment is in?

{quote}

The non-atomicdoBatchOp won't throw IOE....The catch here is weird but it won't be used. We can do a minor refactor for it later.:

{quote}I was trying to see if you could get rid of the cycle looking for Gets in checkAndRowMutate now it has been added to doBatchOp but couldn't figure it; an optimization for another time (smile).

{quote}

ya, another refactor....:(

{quote}On below, I agree.

```
1011 // The atomic ops use the global exception although I feel it is ok to add the exception
```

```
1012 // to each action.
```

{quote}

Oh...Iforgot to remove the comment. After HBASE-19900, our client can handle both of region exception and action exception accurately. Only adding the region exception can reduce the memoryconsumption (it is similar toHBASE-19715).

{quote}We need an issue for the below?

{quote}

Agreed. It would be a nice feature for 2.0.

{quote}We need to make it so it is easy to stand up a RegionServer w/o a Master so you don't have to do the below.

{quote}

The test is necessary for both of branch-1.3 and branch-1.2 as we won't add the such client APIs to them I guess. By contrast, if the impl can be introduced to client APIs for branch-1.4+, we can substitute the real use case for the hacky test.

{quote}+1 to commit.

{quote}

I'm on my travels so I will commit the patch on Monday. Thanks for the reviews [~stack]

15. **stack:** bq. The non-atomic doBatchOp won't throw IOE....The catch here is weird but it won't be used. We can do a minor refactor for it later.

Agree. On commit add in the above line?

Otherwise all good. Makes sense that the test be portable but I should make it so standing up a RegionServer is easy in test; will cut down test times.

+1 on commit Monday.

16. **chia7712:** {quote}On commit add in the above line?

{quote}

roger that

17. **chia7712:** Push this patch to all active branches. Thanks for all reviews.

18. **hudson:** SUCCESS: Integrated in Jenkins build HBase-1.3-IT #348 (See [https://builds.apache.org/job/HBase-1.3-IT/348/])
HBASE-19876 The exception happening in converting pb mutation to (chia7712: rev e36dca55bcd5ad95cd99a19a363)
* (add) hbase-server/src/test/java/org/apache/hadoop/hbase/client/TestMalformedCellFromClient.java
* (edit) hbase-server/src/main/java/org/apache/hadoop/hbase/regionserver/RSRpcServices.java
* (edit) hbase-client/src/main/java/org/apache/hadoop/hbase/protobuf/RequestConverter.java

19. **hudson:** SUCCESS: Integrated in Jenkins build HBase-1.2-IT #1068 (See [https://builds.apache.org/job/HBase-1.2-IT/1068/])
HBASE-19876 The exception happening in converting pb mutation to (chia7712: rev b59c3151dfe41926d242c5c581a2)
* (add) hbase-server/src/test/java/org/apache/hadoop/hbase/client/TestMalformedCellFromClient.java
* (edit) hbase-server/src/main/java/org/apache/hadoop/hbase/regionserver/RSRpcServices.java
* (edit) hbase-client/src/main/java/org/apache/hadoop/hbase/protobuf/RequestConverter.java

20. **hudson:** FAILURE: Integrated in Jenkins build HBase-Trunk__matrix #4579 (See [https://builds.apache.org/job/HBase-Trunk__matrix/4579/])
HBASE-19876 The exception happening in converting pb mutation to (chia7712: rev 2f48fdbb26ff555485b4aa3393d8)

- * (edit) hbase-server/src/main/java/org/apache/hadoop/hbase/regionserver/RSRpcServices.java
- * (edit) hbase-server/src/test/java/org/apache/hadoop/hbase/client/TestMalformedCellFromClient.java
- * (edit) hbase-client/src/main/java/org/apache/hadoop/hbase/shaded/protobuf/RequestConverter.java

Chapter 2

Connected issue HBASE-19900

2.1 Summary

Region-level exception destroy the result of batch

2.2 Description

1) decrease action count repeatedly

If the `AsyncRequestFuture#waitUntilDone` return prematurely, user will get the incorrect results. Or user will be block by `AsyncRequestFuture#waitUntilDone` as the count is never equal with 0.

2) the successive result will be overwritten

3) The failed op is added to `RetriesExhaustedWithDetailsException` repeatedly

`AsyncRequestFutureImpl#receiveMultiAction` process the action-level error first, and then add the region-level exception to each action. Hence, user may get the various exceptions for the same action (row op) from the `RetriesExhaustedWithDetailsException`.

In fact, if both of action-level exception and region-level exception exist, they always have the same context. I'm not sure whether that is what `RetriesExhaustedWithDetailsException` want. As i see it, we shouldn't have the duplicate ops in `RetriesExhaustedWithDetailsException` since that may confuse users if they catch the `RetriesExhaustedWithDetailsException` to check the invalid operations.

2.3 Attachments

1. [HBASE-19900.branch-1.4.v0.patch](#)
2. [HBASE-19900.v0.patch](#)
3. [HBASE-19900.v1.patch](#)
4. [HBASE-19900.v2.patch](#)
5. [HBASE-19900.v2.patch](#)
6. [HBASE-19900.v2.patch](#)

2.4 Comments

1. **chia7712:** The patch merge the check of action-level exception and region-level exception. If the action fail without any exception, we assign the region-level exception to it. If no region-level exception exists, we assign a unknown error instead. In short, each action won't be handled repeatedly so the following issues should be fixed.

decrease action count repeatedly

the successive result will be overwritten

The failed op is added to RetriesExhaustedWithDetailsException repeatedly

The patch also add some debug info to trace the overwrite of result and exception.

2. **chia7712:** All bugs can be easily reproduced by batching the RowMutations. However, only 1.4+ support to do the batch of RowMutations (HBASE-19096). I am not sure whether any real use case can be used to reproduce the bug for branch-1.2 and branch-1.3.
3. **chia7712:** Seems some tests expect that failed ops in RetriesExhaustedWithDetailsException will have the different exceptions.
4. **chia7712:** TestDLSFSHLog is traced by HBASE-19554. Ping for reviews. If no objection against the approach, I will prepare the patch for branch-1.x
5. **yuzhihong@gmail.com:** Looks good overall.

```
1
2 +   Retry canRetry = null;
```

You changed the variable type. Variable name should be changed to reflect this.

cleanServerCache() calls clearCaches(). Better align the method name to be consistent.

```
1
2 +   long getActionsInProgress() {
```

Actions are not returned, only the count. getActionsInProgress -> getNumberOfActionsInProgress

6. **stack:** Skimmed. Looks excellent. +1 Must have been a pain to track down.

7. **chia7712:** v2 patch address [~yuzhihong@gmail.com]’s comment.
8. **yuzhihong@gmail.com:** +1
9. **chia7712:** retry. The failed tests pass locally.
10. **chia7712:** Will commit the fix to all active branches.
11. **chia7712:** Thanks for all reviews.
12. **hudson:** SUCCESS: Integrated in Jenkins build HBase-1.3-IT #347 (See [https://builds.apache.org/job/HBase-1.3-IT/347/])
HBASE-19900 Region-level exception destroy the result of batch (chia7712: rev 32b5f514b1a309375bd51d968ca19ef6
* (add) hbase-client/src/test/java/org/apache/hadoop/hbase/client/TestAsyncProcessWithRegionException.java
* (edit) hbase-client/src/main/java/org/apache/hadoop/hbase/client/AsyncProcess.java
13. **hudson:** SUCCESS: Integrated in Jenkins build HBase-1.2-IT #1067 (See [https://builds.apache.org/job/HBase-1.2-IT/1067/])
HBASE-19900 Region-level exception destroy the result of batch (chia7712: rev ebdffbaf2a090dde5685944604f61a8b3
* (edit) hbase-client/src/main/java/org/apache/hadoop/hbase/client/AsyncProcess.java
* (add) hbase-client/src/test/java/org/apache/hadoop/hbase/client/TestAsyncProcessWithRegionException.java
14. **hudson:** FAILURE: Integrated in Jenkins build HBase-Trunk_matrix #4540 (See [https://builds.apache.org/job/HBase-Trunk_matrix/4540/])
HBASE-19900 Region-level exception destroy the result of batch (chia7712: rev d8b999e6950bc01e8aab8ecce437e710
* (edit) hbase-client/src/main/java/org/apache/hadoop/hbase/client/AsyncRequestFutureImpl.java
* (add) hbase-server/src/test/java/org/apache/hadoop/hbase/client/TestMalformedCellFromClient.java
* (edit) hbase-client/src/main/java/org/apache/hadoop/hbase/client/AsyncBatchRpcRetryingCaller.java
* (add) hbase-client/src/test/java/org/apache/hadoop/hbase/client/TestAsyncProcessWithRegionException.java
15. **xvxinhui:** wonder why you clear the whole cache for a server via invoking cleanServerCache() in this revision, isn’t it too overreacted, any corner case related? [~chia7712]